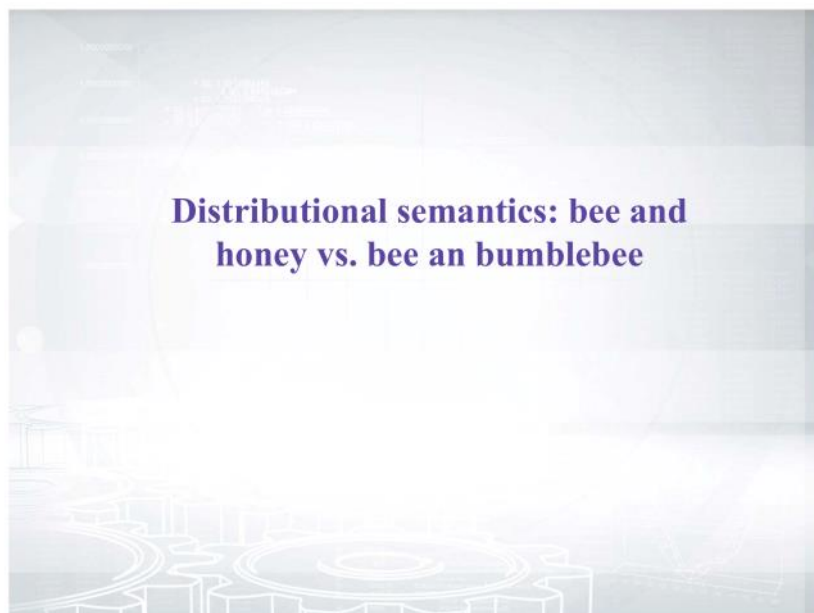





Distributional semantics



Word similarities

- First order co-occurrences
syntagmatic associates / relatedness (bee and honey)
- Second order co-occurrences
paradigmatic parallels / similarity (bee and bumblebee)



Schutze, H., & Pedersen, J. (1993). A vector model for syntagmatic and paradigmatic relatedness. In *Making Sense of Words: Proceedings of the Conference*, pp. 104-113, Oxford, England.

Distributional hypothesis

"You shall know a word by the company it keeps."

— Firth, 1957.

- Use a sliding window of a fixed size
- Compute word co-occurrences n_{uv}

Distributional hypothesis

"You shall know a word by the company it keeps."

— Firth, 1957.

- Use a sliding window of a fixed size
- Compute word co-occurrences n_{uv}
- **Better:** Pointwise Mutual Information:

$$PMI = \log \frac{p(u, v)}{p(u)p(v)} = \log \frac{n_{uv}n}{n_u n_v}$$

Compute a second-order co-occurrence between the words 'bees' and 'honey' (the cosine similarity between their first-order co-occurrence vectors). Use the toy corpus:

These are the wrong sort of bees. Quite the wrong sort. So I should think they would make the wrong sort of honey.

- Let's define a context of a word as three words to the left and three words to the right from the target word, occurred within the same sentence (if there are any).
- For the first-order co-occurrence, let's consider pPMI values (the formula was given on slide 5 of the first video).
Hint: in this question you actually do not need to compute anything... And the answer would be the same for any type of first-order co-occurrence.

From <https://www.coursera.org/learn/deep-neural-network/exercises/week4/word-and-sentence-embeddings>

whether the words
randomly occur
or they are
related.

	bee	honey
wrong sort of	1	1
quite the	1	1
so I should think they	1	1
would make	1	1
the	1	1

if never co-occur - $\frac{p(u, v)}{p(u)p(v)} \approx 0$

$\Rightarrow \text{pmi} \rightarrow -\infty$

so, take max $(\log \frac{n_{uv}n}{n_u n_v}, 0)$

Distributional hypothesis

"You shall know a word by the company it keeps."

— Firth, 1957.

- Use a sliding window of a fixed size
- Compute word co-occurrences n_{uv}
- **Better:** Pointwise Mutual Information:

$$PMI = \log \frac{p(u, v)}{p(u)p(v)} = \log \frac{n_{uv}n}{n_u n_v}$$

- **Even better:** positive Pointwise Mutual Information:

$$pPMI = \max(0, PMI)$$

Any problems here?


- First order co-occurrences
syntagmatic associates / relatedness (bee and honey)
- Second order co-occurrences
paradigmatic parallels / similarity (bee and bumblebee)



Schutze, H., & Pedersen, J. (1993). A vector model for syntagmatic and paradigmatic relatedness. In *Making Sense of Words: Proceedings of the Conference*, pp. 104-113, Oxford, England.


Vector Space Models of Semantics

- **Input:** word-word co-occurrences (counts, PMI, ...)
- **Method:** dimensionality reduction (SVD, ...)
- **Output:** similarity between vector representations of words



Turnay, P.D., Pantel, P.: from Frequency to Meaning: Vector Space Models of Semantics, 2010.


What is a context?



WORD	CONTEXTS
australian	scientist/ <u>amod</u> ⁻¹
scientist	australian/ <u>amod</u> , discovers/ <u>nsubj</u> ⁻¹
discovers	scientist/ <u>nsubj</u> , star/ <u>dobj</u> , telescope/ <u>prep_with</u>
star	discovers/ <u>dobj</u> ⁻¹
telescope	discovers/ <u>prep_with</u> ⁻¹

Omer Levy, Yoav Goldberg, Dependency-Based Word Embeddings, ACL-2014.

Syntax helps to understand what is local context and what is not



matrix factorization




2. Choose correct statements about Singular Value Decomposition (SVD), an important notion from the linear algebra. Feel free to consult any additional resource like [wiki](#) if needed.

- ☒ Truncated SVD is the best rank k approximation of the original matrix in terms of Frobenius norm.
- ☒ Any rectangular matrix with real entries has a singular value decomposition.
- ☐ Singular values can be negative.
- ☒ Squares of singular values of a matrix X are eigenvalues of $X^T X$ (or XX^T).
- ☐ Singular values of a rectangular matrix are its eigenvalues.
- ☒ Singular values decomposition is not unique (for example, the zero matrix can be decomposed in infinitely many ways).


Truncated SVD

Keep only first k components: $\hat{X}_k = U_k \Sigma_k V_k^T$



Truncated SVD

Keep only first k components: $\hat{X}_k = U_k \Sigma_k V_k^T$

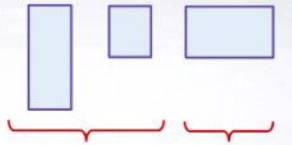


It's the best approximation of rank k in terms of Frobenius norm:

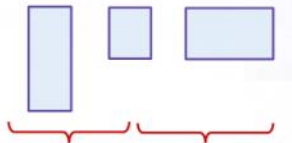
$$\|X - \hat{X}\|_F = \sqrt{\sum_{i=1}^n \sum_{j=1}^m (x_{ij} - \hat{x}_{ij})^2}$$

How do we use it?

Option 1:


$$\Phi = U_k \Sigma_k \quad \Theta = V_k^T$$


Option 2:


$$\Phi = U_k \sqrt{\Sigma_k} \quad \Theta = \sqrt{\Sigma_k} V_k^T$$

makes two matrices

Vector Space Models of Semantics

- **Input:** word-word co-occurrences (counts, PMI, ...)
- **Method:** dimensionality reduction (SVD, ...)
- **Output:** similarity between vector representations of words




Turnay, P.D., Pantel, P.: from Frequency to Meaning: Vector Space Models of Semantics, 2010.

Weighted squared loss: GloVe

Fill X with $\log n_{uv}$ and try another objective:

$$\sum_{u \in W} \sum_{v \in W} f(n_{uv}) (\langle \phi_u, \theta_v \rangle + b_u + b'_v - \log n_{uv})^2 \rightarrow \min_{\phi_u, \theta_v, b_u, b'_v}$$



Pennington et. al. GloVe: Global Vectors for Word Representation, 2014.

matrix factorization

original matrices. log weighted square loss.

caps at 1, otherwise too frequent words get too much weights.

Word prediction: skip-gram model

Predict context words given a focus word:

$$p(w_{i-h}, \dots, w_{i+h} | w_i) = \prod_{-h \leq k \leq h, k \neq 0} p(w_{i+k} | w_i)$$

Model each probability with a softmax:

$$\underline{p(u|v)} = \frac{\exp \langle \phi_u, \theta_v \rangle}{\sum_{u' \in W} \exp \langle \phi_{u'}, \theta_v \rangle}$$


Still two matrices of parameters.

Training this model.
SGD.

How do we train the model?

Log-likelihood maximization:

$$\mathcal{L} = \sum_{u \in W} \sum_{v \in W} n_{uv} \log p(u|v)$$

 *word co-occurrence*

Method:

- SGD, online by word pairs in the corpus

Problem:

- *softmax* over vocabulary is slow!

Skip-gram Negative Sampling (SGNS)

Instead of predicting a word for another word,
predict “yes” or “no” for word pairs:

$$\sum_{u \in W} \sum_{v \in W} n_{uv} \log \sigma(\langle \phi_u, \theta_v \rangle) +$$

$$k \mathbb{E}_{\bar{v}} \log \sigma(-\langle \phi_u, \theta_{\bar{v}} \rangle) \rightarrow \max_{\phi_u, \theta_v}$$

- Use **positive examples** from data: v co-occurred with u
- Sample **negative examples**: k random \bar{v} from the vocabulary

Train with SGD to find two matrices of parameters (as usual).


sample for every u -word

sigmoid function

SGNS as implicit matrix factorization

SGNS objective is maximized when $\langle \phi_u, \theta_v \rangle$ is equal to shifted Pointwise Mutual Information:

$$\text{sPMI} = \log \frac{n_{uv}n}{n_u n_v} - \log k$$

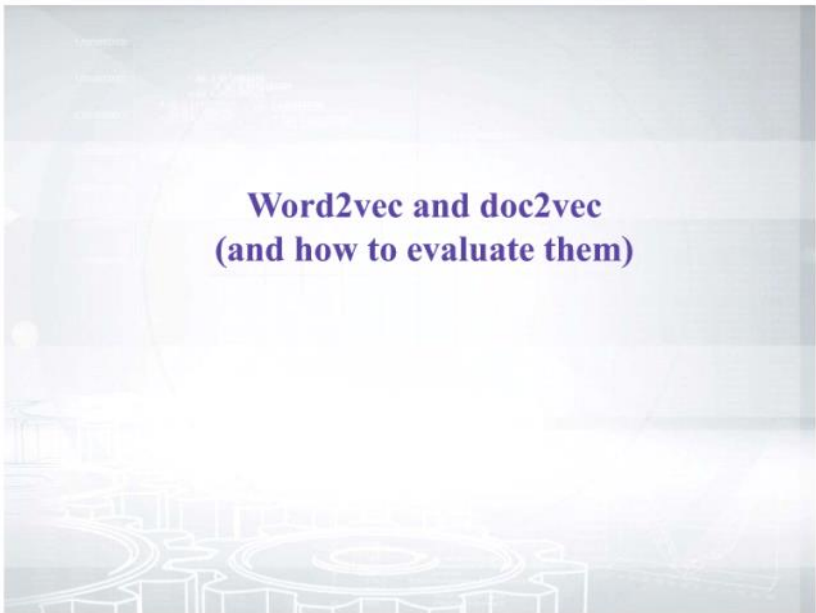


Levy and Goldberg. Neural Word Embedding as Implicit Matrix Factorization, 2014.

even we don't think about matrix factorization, it's still implicitly matrix factorization.



word2vec and doc2vec

The background of the slide features a faint, stylized illustration. It includes a large globe in the upper half and a series of interlocking gears in the lower half, suggesting a theme of global technology or engineering. The text is centered over this background.

Word2vec and doc2vec (and how to evaluate them)

Word2vec

Two architectures:

- CBOW (Continuous Bag-of-words):

$$p(w_i | w_{i-h}, \dots, w_{i+h})$$

context \rightarrow word of interest

- Continuous Skip-gram:

$$p(w_{i-h}, \dots, w_{i+h} | w_i)$$

word of interest \rightarrow context

Two ways to avoid softmax:

- Negative sampling
- Hierarchical softmax

Open-source and fast: code.google.com/archive/p/word2vec/

Evaluation: word similarities

How do we test that similar words have similar vectors?

- Linguists know a lot about what is “similar”.
- We can use human judgements for word pairs.
- Compare Spearman's correlation between two lists:

Linguists' table.

tiger	tiger	10.00
media	radio	7.42
tiger	cat	7.37
train	car	6.31
...


tiger	tiger	$\cos(\phi_u, \phi_v)$
media	radio	...
tiger	cat	...
train	car	...
...

not straightforward task
Not the best way.

Evaluation: word analogies

- In cognitive science well known as *relational similarity* (vs. *attributional similarity*).
- $a : a'$ is as $b : b'$ (man : woman is as king : ?)

$$\cos(b - a + a', x) \rightarrow \max_x$$



Gentner, D. Structure-mapping: A theoretical framework for analogy. Cognitive Science, 1983.
Mikolov et. al. Linguistic Regularities in Continuous Space Word Representations, 2013.

relational similarity.

~~we~~ before. we talk about attributional similarity.

Word similarity task performance

- For word similarity task, count-based methods (PPMI, SVD) perform on par with predictive methods (GloVe, SGNS).

win	Method	WordSim Similarity	WordSim Relatedness	Bruni et al. MEN	Radinsky et al. M. Turk
2	PPMI	.732	.699	.744	.654
	SVD	.772	.671	.777	.647
	SGNS	.789	.675	.773	.661
	GloVe	.720	.605	.728	.606
5	PPMI	.732	.706	.738	.668
	SVD	.764	.679	.776	.639
	SGNS	.772	.690	.772	.663
	GloVe	.745	.617	.746	.631

win is the width of the window for co-occurrences collection.

Levy et. al. Improving distributional similarity with lessons learned from word embeddings, 2015.

what's GloVe

[Picture on page "Word and sentence embedding"](#)

→ matrix factorization of log-counts
wrt weighted squared loss

Word analogy task performance

- Word analogy task is solved with 70% average accuracy.

win	Method	Google		MSR	
		Add	Mul	Add	Mul
2	PPMI	.552	.677	.306	.535
	SVD	.554	.591	.408	.468
	SGNS	.676	.689	.617	.644
	GloVe	.649	.666	.540	.591
5	PPMI	.518	.649	.277	.467
	SVD	.532	.569	.369	.424
	SGNS	.692	.714	.605	.645
	GloVe	.700	.712	.541	.599

Add is the way of analogy solving that we discussed. Mul is a modification.

Levy et. al. Improving distributional similarity with lessons learned from word embeddings, 2015.



Paragraph2vec aka doc2vec

And the only reason for being a bee that I know of is making honey.

contexts

↑
focus
word

contexts

DM (Distributed Memory):

$$p(w_i | w_{i-h}, \dots, w_{i+h}, d)$$

DBOW (Distributed Bag Of Words):

$$p(w_{i-h}, \dots, w_{i+h} | d)$$

Condition not on
word, but on document

Evaluation: document similarities

How do we test that *similar documents* have *similar vectors*?

- ArXiv triplets: paper A, similar paper B, dissimilar paper C
- Measure the accuracy of guessing the dissimilar paper

http://arxiv.org/pdf/1206.5743	http://arxiv.org/pdf/cond-mat/0403258	http://arxiv.org/pdf/1408.0189
http://arxiv.org/pdf/1209.0268	http://arxiv.org/pdf/1307.7598	http://arxiv.org/pdf/math/0504051
http://arxiv.org/pdf/hep-ph/9908436	http://arxiv.org/pdf/nucl-th/9707019	http://arxiv.org/pdf/1112.3014
http://arxiv.org/pdf/1111.2905	http://arxiv.org/pdf/1303.2538	http://arxiv.org/pdf/1109.1922
http://arxiv.org/pdf/nucl-ex/0112013	http://arxiv.org/pdf/physics/9704013	http://arxiv.org/pdf/1408.4595
http://arxiv.org/pdf/0709.3419	http://arxiv.org/pdf/quant-ph/0611134	http://arxiv.org/pdf/0902.0616
http://arxiv.org/pdf/hep-th/9609148	http://arxiv.org/pdf/solv-int/9710009	http://arxiv.org/pdf/astro-ph/0508060



Andrew Dai, Cristopher Olah, Quoc Le. Document Embedding with Paragraph Vectors, CoRR, 2015.



Evaluation: document similarities

Integral formula of Minkowski type and new characterization of the Wulff shape

Yijun He ^{*} Haiahong Li [†]

Abstract

Given a positive function F on S^n which satisfies a convexity condition, we introduce the r -th anisotropic mean curvature M_r for hypersurfaces in \mathbb{R}^{n+1} which is a generalization of the usual r -th mean curvature H_r . We get integral formulas of Minkowski type for compact hypersurfaces in \mathbb{R}^{n+1} . We give some new characterizations of the Wulff shape by use of our integral formulas of Minkowski type, in case $F = 1$ which reduces to some well-known results.

2000 Mathematics Subject Classification: Primary 53C42, 53A30; Secondary 53B25.

Key words and phrases: Wulff shape, F -Weingarten operator, anisotropic principal curvature, r -th anisotropic mean curvature, integral formula of Minkowski type.

[xiv.org/pdf/1408.0189](https://arxiv.org/pdf/1408.0189)
[xiv.org/pdf/math/0504051](https://arxiv.org/pdf/math/0504051)
[xiv.org/pdf/1112.3014](https://arxiv.org/pdf/1112.3014)
[xiv.org/pdf/1109.1922](https://arxiv.org/pdf/1109.1922)
[xiv.org/pdf/1408.4595](https://arxiv.org/pdf/1408.4595)
[xiv.org/pdf/0902.0616](https://arxiv.org/pdf/0902.0616)
[xiv.org/pdf/astro-ph/0508060](https://arxiv.org/pdf/astro-ph/0508060)

Evaluation: document similarities

Integral formula of Minkowski type and new
characterization of the Wulff shape

**COMPLEX CURVES IN ALMOST-COMPLEX
MANIFOLDS AND MEROMORPHIC HULLS**

Sergei IVASHKOVICH – Vsevolod SHEVCHISHIN

Preface

Chapter I. Local Properties of Complex Curves.

Lecture 1. Complex Curves in Almost-Complex Manifolds. ... pp. 1–12

1.1. Almost-Complex Manifolds, Hermitian Metrics, As-
sociated (1,1)-Forms. 1.2. Existence of Calibrating and
Tame Structures. 1.3. Almost-Complex Submanifold,
Complex Curves, Energy and Area. 1.4. Symplectic Sur-
faces. 1.5. Adjunction Formula for Immersed Symplectic
Surfaces.

Evaluation: document similarities

Accepted for publication in *Solar Physics*, waiting for the authoritative version and a DOI which will be available at <http://www.springerlink.com/content/0038-0938>

Time-dependent Stochastic Modeling of Solar Active Region Energy

M. Kanazir and M. S. Wheatland¹

Received: 7 July 2010 / Accepted: 31 July 2010 / Published online:

Abstract A time-dependent model for the energy of a flaring solar active region

Lecture 1. Complex Curves in Almost-Complex Manifolds. ... pp. 1-12

1.1. Almost-Complex Manifolds, Hermitian Metrics, Associated (1,1)-Forms. 1.2. Existence of Calibrating and Tame Structures. 1.3. Almost-Complex Submanifold, Complex Curves, Energy and Area. 1.4. Symplectic Surfaces. 1.5. Adjunction Formula for Immersed Symplectic Surfaces.

2000 Ma
53H25.

Key word
curvature,

Resume

Methods:

- *word2vec*: SGNS, CBOW, ...
- *doc2vec*: DBOW, DM, ...
- Python library for both: <https://radimrehurek.com/gensim/>

Evaluation:

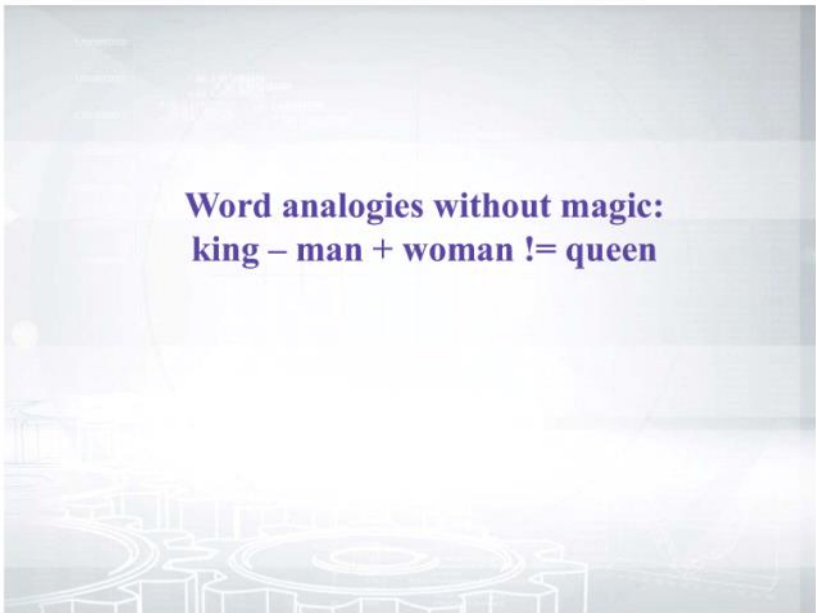
- Word similarity and analogy
- Document similarity
- *Interpretability of the components*
- *Geometry of the embeddings space*

Count-based and predictive approaches are not so different!

count-based method
not much different
from
predictive-based method!



word analogy

The background of the slide is a faded, abstract image. It features a cityscape with a prominent bridge, possibly the Golden Gate Bridge, and a series of interlocking gears in the foreground, suggesting a theme of technology, engineering, or complex systems.

**Word analogies without magic:
king – man + woman != queen**

A magical property of word2vec

And the only reason for being a **bee** that I know of is making honey.

contexts focus word contexts

Learn word vectors by predicting their contexts (or vice-versa).

Obtain vectors that solve word analogies:

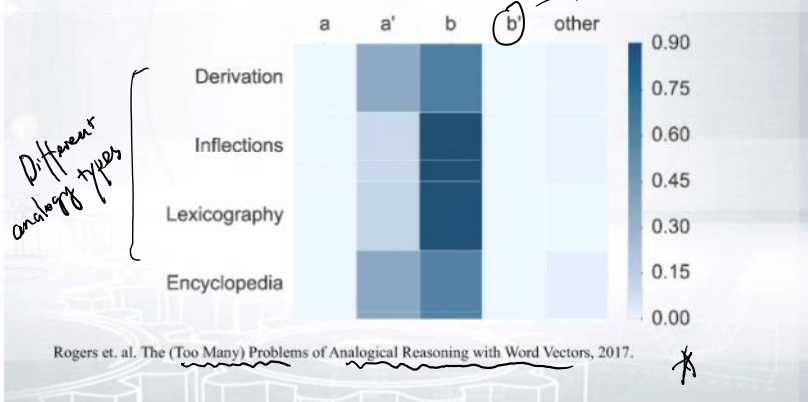
- king – man + woman = queen
- Moscow – Russia + France = Paris

Demo: rare-technologies.com/word2vec-tutorial/

Closer look into analogy task

$$\cos(b - a + a', x) \rightarrow \max_{x \notin \{a, a', b\}}$$

king - man + woman = king:




never close to b' — queen the target

If we do not exclude b.
we end up with b.
是否还排除 b?

The closest neighbor of b is good enough?

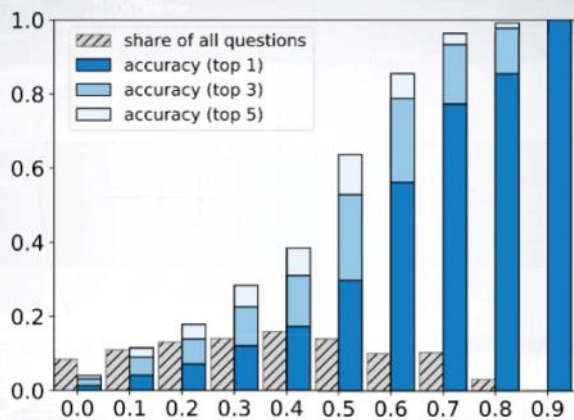
For the plural noun category in the Google test set:
70% accuracy by taking the closest neighbor of the vector b !



Linzen. Issues in evaluating semantic spaces using word analogies, 2016.

不管减掉了谁。
找 b 的 ~~the~~ closest vector
就好了。
high accuracy

Good accuracy when b and b' are close



Buckets are based on similarity between b and b'

Split analogy examples
by the similarity of b & b'

⇒ works really nice when b and b'
are similar.
works poorly when b & b'
are different (complicated tasks).

BATS dataset

Inflectional morphology:

- *student:students, strong:stronger, follow:following, ...*

Derivational morphology:

- *bake:baker, edit:editable, home:homeless, ...*

Lexicographic semantics:

- **Hypernyms, meronyms** : peach:fruit, sea:water, player:team, ...
- **Antonyms, synonyms**: up:down, clean:dirty, cry:scream, ...

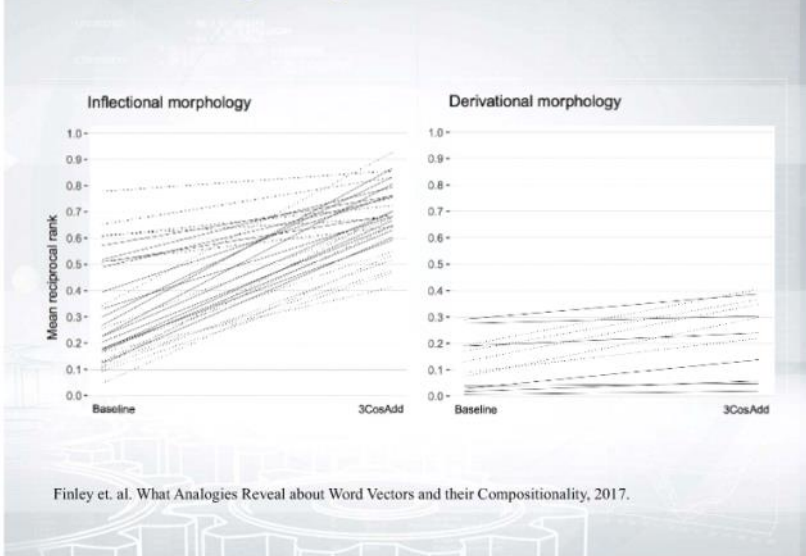
Encyclopedic semantics:

- **Animals**: cat:kitten, dog:bark, ...
- **Geography**: Athens:Greece, Peru:Spanish, ...
- **People**: Lincoln:president, Lincoln:American, ...
- **Other**: blood:red, actor:actress, ...

让A2做类比还是不给力吧.

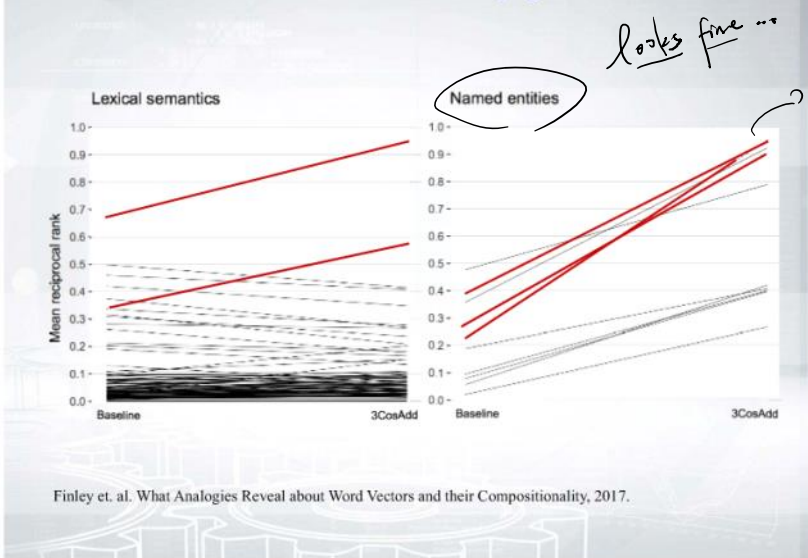
Compare to baseline:
just use the closest

Performance by categories



horizontal
no difference
high difference

Genders and counties are cherry-picks



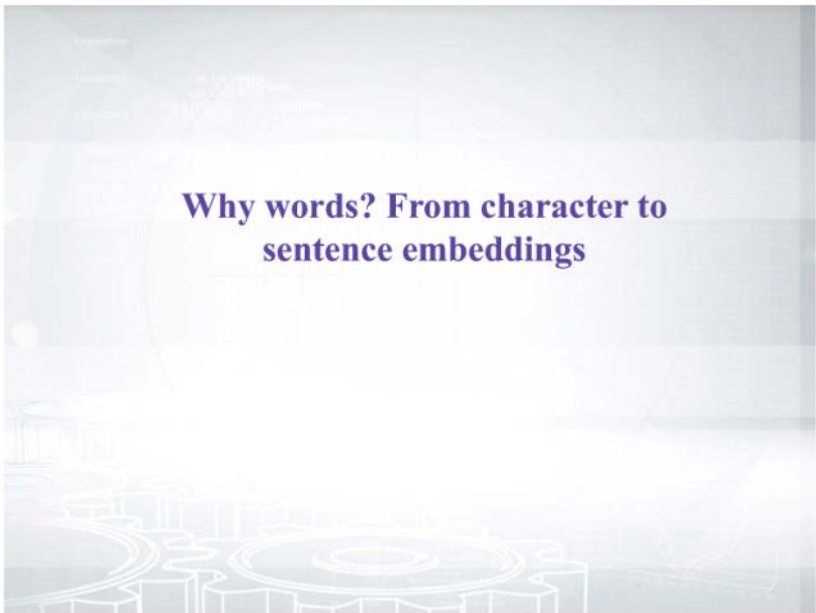
Resume

- *word2vec* works fine for word similarities
- But there are many questions with word analogies
- Be careful about hype!

Does not solve all analogy tasks.
Nice area for future research



why words

The background of the slide features a complex, abstract design. It includes a series of interlocking gears at the bottom, rendered in a light blue and white color scheme. Above the gears, there are faint, overlapping lines and shapes that suggest a data visualization or a network diagram. The overall color palette is dominated by light blues, greys, and whites, with a subtle gradient effect.

Why words? From character to sentence embeddings


Morphology can help

Problems:

- Languages with rich morphology
- Low frequency words, OOV words

morphology

Need linguistic rules



Use morphology to improve word embeddings

I. Vulic et. al. Morph-fitting: Fine-Tuning Word Vector Spaces with Simple Language-Specific Rules, 2017.

FastText

Represent a word as a bag of character n-grams, e.g. for $n = 3$:

$G_{where} : \text{_wh, whe, her, ere, re_, _where_}$

Model a word vector as a sum of sub-word vectors:

SGNS:

$$\text{sim}(u, v) = \langle \phi_u, \theta_v \rangle$$

FastText:

$$\text{sim}(u, v) = \sum_{g \in G_v} \langle \phi_u, \theta_g \rangle$$

Code and pre-trained embeddings: <https://fasttext.cc/>

P. Bojanowsky et al. Enriching Word Vectors with Subword Information, 2016.

FastText

word vector as sum of subword vectors



Sent2vec

> not nice approach.

First ideas:

- Average pre-trained word vectors (*word2vec*, *GloVe*, etc).
- Maybe use TF-IDF weights for averaging. *use TF-IDF as weights*

Sent2vec:

- Learn sentence embedding as a sum of sub-sentence units: *Sub-sentence units.*

$$\text{sim}(u, s) = \frac{1}{|G_s|} \sum_{g \in G_s} \langle \phi_u, \theta_g \rangle$$

where G_s is a set of word n-grams for the sentence s .

Code and embeddings: <https://github.com/epfml/sent2vec>

Pagliardini et. al. Unsupervised Learning of Sentence Embeddings using Compositional n-Gram Features, 2017.

★ Similar to FastText
 #. Just different units -
 word n-grams to represent
 sentence.

Have you realized what's the crucial difference between averaging word2vec vectors and sent2vec?

- ✓ sent2vec represents a sentence not as an average of words, but as an average of n-grams (SUB-SENTENCE UNITS)

- sent2vec represents sentences as an average of words DURING TRAINING, while averaging word2vec vectors is a postprocessing step

When $n=1$ then it's the same as averaging word-to-vec

This should not be selected

Even though sent2vec is still based on word (or n-gram) vectors, it tunes them during training to fit word-in-sentence co-occurrence data. This is not the same as fitting word-word co-occurrence data (as word2vec does) and can crucially improve the final representations of sentences.

StarSpace

General framework:

entities (e.g. sentences) and features (e.g. words)

Lot's of applications:

- Text classification, e.g. sentiment analysis
- Ranking, e.g. ranking web documents given a query
- Collaborative filtering-based recommendation
- Content-based recommendation
- Embedding graphs, e.g. Facebook

- Collaborative filtering-based recommendation
- Content-based recommendation
- Embedding graphs, e.g. Freebase
- Learning word, sentence or document embeddings

Code and tutorials: github.com/facebookresearch/Starspace *

Wu et. al. StarSpace: Embed All The Things! 2017

Embed All the things!

StarSpace

Mode 3 (sentence embeddings):

Use case: learn pairwise similarity from collections of similar objects, e.g. sentence similarity. *

Data format: each line is a collection of similar sentences:

sent1_word1 sent1_word2 ... <tab> sent2_word1 sent2_word2 ...

Training:

- Each sentence is represented as a bag of features (words or n-grams). They are embedded to predict sentence similarity
- Similar sentence pairs are taken from the collections
- Dissimilar sentence pairs are sampled at random

Deep learning?

The most popular options:

- Recurrent Neural Networks (sequence modelling)
- Convolutional Neural Networks (much faster)
- Recursive Neural Networks (use hierarchical structure)

Linguistic structure is back:


- Morphology can help to build word embeddings
- Recursive Neural Networks, Tree-LSTMs, DAG-LSTMs, etc. use syntax, span annotations, co-reference links...

* use language to build hierarchical representation

Skip-thought vectors

- Predict next and previous sentences in text
- RNN encoder-decoder architecture

Thought vector



Kiros et. al. Skip-Thought Vectors, 2015, <https://github.com/ryankiros/skip-thoughts>.

encoder-decoder architecture
不是用来翻译。
而是 predict next sentence.
也是有一定的能力的。